

Workshop 1: Language and NumPy basics

FIE463: Numerical Methods in Macroeconomics and Finance using Python

Richard Foltyn
NHH Norwegian School of Economics

January 16, 2025

See GitHub repository for notebooks and data:

<https://github.com/richardfoltyn/FIE463-V25>

1 Type conversions

In the lecture, we discussed the basic built-in data types: integers, floating-point numbers, booleans, and strings. Python allows us to convert one type to another using the following functions:

- `int()` converts its argument to an integer.
- `float()` converts its argument to a floating-point number.
- `bool()` converts its argument to a boolean.
- `str()` converts its argument to a string.

These conversions mostly work in an intuitive fashion, with some exceptions.

1. Define a string variable `s` with the value `'1.1'`. Convert this variable to integer, float, and boolean.
2. Define the string variables `s1`, `s2`, and `s3` with values `'True'`, `'False'`, and `''` (empty string), respectively. Convert each of these to boolean. Can you guess the conversion rule?
3. Define a floating-point variable `x` with the value `0.9`. Convert this variable to integer, boolean, and string.
4. Define the integer variables `i1` and `i2` with values `0` and `2`, respectively. Convert each of these variables to boolean.
5. Define the boolean variables `b1` and `b2` with values `True` and `False`, respectively. Convert each of them to integer.
6. NumPy arrays cannot be converted using `int()`, `float()`, etc. Instead, we have to use the method `astype()` and pass the desired data type (e.g., `int`, `float`, `bool`) as an argument. Create a NumPy array called `arr` with elements `[0.0, 0.5, 1.0]` and convert it to integer and boolean type.

Solution.

Part (1)

```
[1]: # Define the string variable  
s = '1.1'
```

```
[2]: # Attempt to convert to integer: this fails because string contains decimal digits  
int(s)
```

```
ValueError: invalid literal for int() with base 10: '1.1'
```

```
[3]: # Convert to float  
float(s)
```

```
[3]: 1.1
```

```
[4]: # Convert to boolean: this returns True because any non-empty string is considered True  
bool(s)
```

```
[4]: True
```

Part (2)

```
[5]: # Define string variables  
s1 = 'True'  
s2 = 'False'  
s3 = ''
```

```
[6]: # Convert 'True' to boolean  
bool(s1)
```

```
[6]: True
```

```
[7]: # Convert 'False' to boolean  
bool(s2)
```

```
[7]: True
```

```
[8]: # Convert '' to boolean  
bool(s3)
```

```
[8]: False
```

Part (3)

```
[9]: # Define floating-point number  
x = 0.9
```

```
[10]: # Convert to integer: The decimal digits are truncated, NOT rounded!  
int(x)
```

```
[10]: 0
```

```
[11]: # Convert to boolean: Any numeric value other than 0 is interpreted as True  
bool(x)
```

[11]: True

```
[12]: # Convert to string  
      str(x)
```

[12]: '0.9'

Part (4)

```
[13]: # Define integer variables  
      i1 = 0  
      i2 = 2
```

```
[14]: # Convert 0 to boolean: 0 numerical values are interpreted as False  
      bool(i1)
```

[14]: False

```
[15]: # Convert 2 to boolean: any non-zero numerical value is interpreted as True  
      bool(i2)
```

[15]: True

Part (5)

```
[16]: # Define boolean variables  
      b1 = True  
      b2 = False
```

```
[17]: # Convert True to integer  
      int(b1)
```

[17]: 1

```
[18]: # Convert False to integer  
      int(b2)
```

[18]: 0

Part (6)

```
[19]: # Import numpy  
      import numpy as np  
  
      # Create array of floating-point numbers  
      arr = np.array([0.0, 0.5, 1.0])  
      arr
```

[19]: array([0. , 0.5, 1.])

```
[20]: # Convert to integer array: note that values are again truncated, NOT rounded.  
      arr.astype(int)
```

[20]: array([0, 0, 1])

```
[21]: # Convert to boolean array
arr.astype(bool)
```

```
[21]: array([False,  True,  True])
```

2 Working with strings

Strings in Python are full-fledged objects, i.e., they contain both the character data as well as additional functionality implemented via functions or so-called *methods*. The official [documentation](#) provides a comprehensive list of these methods. For our purposes, the most important are:

- `str.lower()` and `str.upper()` convert the string to lower or upper case, respectively.
- `str.strip()` removes any leading or trailing whitespace characters from a string.
- `str.count()` returns the number of occurrences of a substring within a string.
- `str.startswith()` and `str.endswith()` check whether a string starts or ends with a given substring.

Moreover, strings are also sequences, and as such support indexing in the same way as lists or tuples.

Create a string variable with the value

```
s = '  NHH Norwegian School of Economics  '
```

and perform the following tasks:

1. Strip and surrounding spaces from the string using `strip()`.
2. Count the number of 'H' in the string.
3. Modify your code so that it is case-insensitive, i.e., both instances of 'h' and 'H' are counted.
4. Reverse the string, i.e., the last character should come first, and so on.
5. Create a new string which contains every 2nd letter from the original.
6. Select the last character from this new string using at least two different methods.

Solution.

Part (1)

```
[22]: # Define the string
s = '  NHH Norwegian School of Economics  '

# strip leading & trailing whitespace
s = s.strip()
s
```

```
[22]: 'NHH Norwegian School of Economics'
```

Part (2)

```
[23]: # Count the number of H in string
      s.count('H')
```

```
[23]: 2
```

Part (3)

```
[24]: # Count number of H, ignore case:
      # We first convert the string to lower case, then look for number of h
      s.lower().count('h')
```

```
[24]: 3
```

Part (4)

```
[25]: # Reverse the string
      s[::-1]
```

```
[25]: 'scimonocE fo loohcS naigewroN HHN'
```

Recall that the slice `start:stop:step` can be used to index sequences *and* strings. We use the default values for `start` and `stop` (so they can be omitted), but set the `step` to `-1` to reverse the order.

Part (5)

```
[26]: s2 = s[::-2]
      s2
```

```
[26]: 'NHNreinSho fEoois'
```

Part (6)

```
[27]: # Select last element using -1
      s2[-1]
```

```
[27]: 's'
```

```
[28]: # Select last element using len()
      s2[len(s2) - 1]
```

```
[28]: 's'
```

3 Summing lists and arrays

In this exercise, we investigate an additional difference between built-in lists and NumPy arrays: performance. You are asked to investigate performance differences for different implementations of the `sum()` function.

1. Create a list `lst` and a NumPy array `arr`, each of them containing the sequence of ten values 0, 1, 2, ..., 9.

Hint: You can use the list constructor `list()` and combine it with the `range()` function which returns an object representing a range of integers.

Hint: You should create the NumPy array using `np.arange()`.

2. We want to compute the sum of integers contained in `lst` and `arr`. Use the built-in function `sum()` to sum elements of a list. For the NumPy array, use the NumPy function `np.sum()`.
3. You are interested in benchmarking which summing function is faster. Repeat the steps from above, but use the cell magic `%timeit` to time the execution of a statement as follows:

`%timeit statement`

4. Recreate the list and array to contain 100 integers starting from 0, and rerun the benchmark.
5. Recreate the list and array to contain 10,000 integers starting from 0, and rerun the benchmark.

What do you conclude about the relative performance of built-in lists vs. NumPy arrays?

Solution.

Part (1)

```
[29]: # create list with 10 elements 0,1,...,9
      lst = list(range(10))
```

```
[30]: import numpy as np
      # create array with 10 elements 0,1,...,9
      arr = np.arange(10)
```

Part (2)

```
[31]: # sum list using the built-in function sum()
      sum(lst)
```

```
[31]: 45
```

```
[32]: # sum the NumPy array using NumPy's sum() function
      np.sum(arr)
```

```
[32]: np.int64(45)
```

Part (3)

```
[33]: # benchmark summing list using built-in sum()
      %timeit sum(lst)
```

60.3 ns ± 1.01 ns per loop (mean ± std. dev. of 7 runs, 10,000,000 loops each)

```
[34]: # Benchmark summing array using NumPy's sum()
      %timeit np.sum(arr)
```

1.56 µs ± 30.9 ns per loop (mean ± std. dev. of 7 runs, 1,000,000 loops each)

As you can see, for a short list the built-in `sum()` was faster by a factor of about 25 (the exact difference varies depending on your hardware and platform).

Part (4)

```
[35]: # Recreate list and array to contain 100 integers starting at 0
      N = 100
      lst = list(range(N))
      arr = np.arange(N)
```

```
[36]: # benchmark built-in sum() with 100 elements
      %timeit sum(lst)
```

336 ns \pm 1.95 ns per loop (mean \pm std. dev. of 7 runs, 1,000,000 loops each)

```
[37]: # benchmark NumPy's sum() with 100 elements
      %timeit np.sum(arr)
```

1.53 μ s \pm 14 ns per loop (mean \pm std. dev. of 7 runs, 1,000,000 loops each)

For 100 elements, the built-in `sum()` is still faster, but only by a factor of 4 (again, the exact values depend on your platform). Note that the execution time for `np.sum()` remained almost unchanged, which suggest that the function call has a high fixed cost, but scales much better with the number of elements to be summed.

Part (5)

```
[38]: # Recreate list and array to contain 10,000 integers starting at 0
      N = 10000
      lst = list(range(N))
      arr = np.arange(N)
```

```
[39]: # benchmark built-in sum() with 10000 elements
      %timeit sum(lst)
```

49.9 μ s \pm 218 ns per loop (mean \pm std. dev. of 7 runs, 10,000 loops each)

```
[40]: # benchmark NumPy's sum() with 10000 elements
      %timeit np.sum(arr)
```

2.75 μ s \pm 36.6 ns per loop (mean \pm std. dev. of 7 runs, 100,000 loops each)

Lastly, for 10,000 elements `np.sum()` is substantially faster by a factor about 20. You should conclude that for large arrays, you can expect much better performance from NumPy's functions, but this may not be true for small data sets.
